

Problem Statement:

Task Objective

Design and implement two distinct news article recommendation algorithms that leverage a user's past reading behavior. The algorithms should consider:

- User's reading history
- User's expressed interests
- Popularity of news articles
- Relevance of articles to the user's current location

The developed algorithms must be capable of real-time recommendations and suitable for deployment in a production environment.

Data:

- Device Data : gs://nis-interview-task-data/devices
- Event Data : gs://nis-interview-task-data/event
- Testing Content : gs://nis-interview-task-data/testing_content
- Training Content : gs://nis-interview-task-data/training_content

Exploratory Data Analysis:

- **User**

The dataset contains **10,400 user records** with **11 columns** capturing device, platform, location, language, and activity details.

- **Identifiers & device info:**

Every user has a unique **deviceid**. All users have Android phones. Device **model** is available for most users (99.5%) but missing for a small fraction (~56 records).

- **Network info:**

The **networkType** column is almost complete (only 48 missing).


- **Location fields:**

district is **mostly missing** (only 21 non-null values out of 10,400), so it's not reliable. In contrast, **lastknownsubadminarea** (user's city) is much better populated, available for ~91% of users.

- **Language preference:**

The **language_selected** is not useful as the only language available is **en**

User Table EDA Dashboard



- **Event**

- ◆ Event Distribution

- **timespent-front** dominates (~3.6M events).
- Other events (**timespent-back**, bookmarks, shares, searches) are negligible in comparison.

- ◆ Time Spent

- Majority of events have <10 sec time spent.
- **timespent-back** shows the highest median and variability.
- Other event types contribute little to overall time.

- ◆ Card View Position

- Most engagement happens in the top 20 card positions.
- Very steep drop-off after position 50; minimal activity beyond 100.

- ◆ Temporal Trends

- Hourly: Peak activity between 9 AM – 9 PM, low activity overnight (12–4 AM).
- Weekly: Highest on Tuesday, lowest on weekends (Sat & Sun).

- ◆ Relevancy Feedback

- Overwhelmingly labeled as “unknown” (~3.5M).
- Very few **green**, **yellow**, **red** labels → feedback data is sparse.

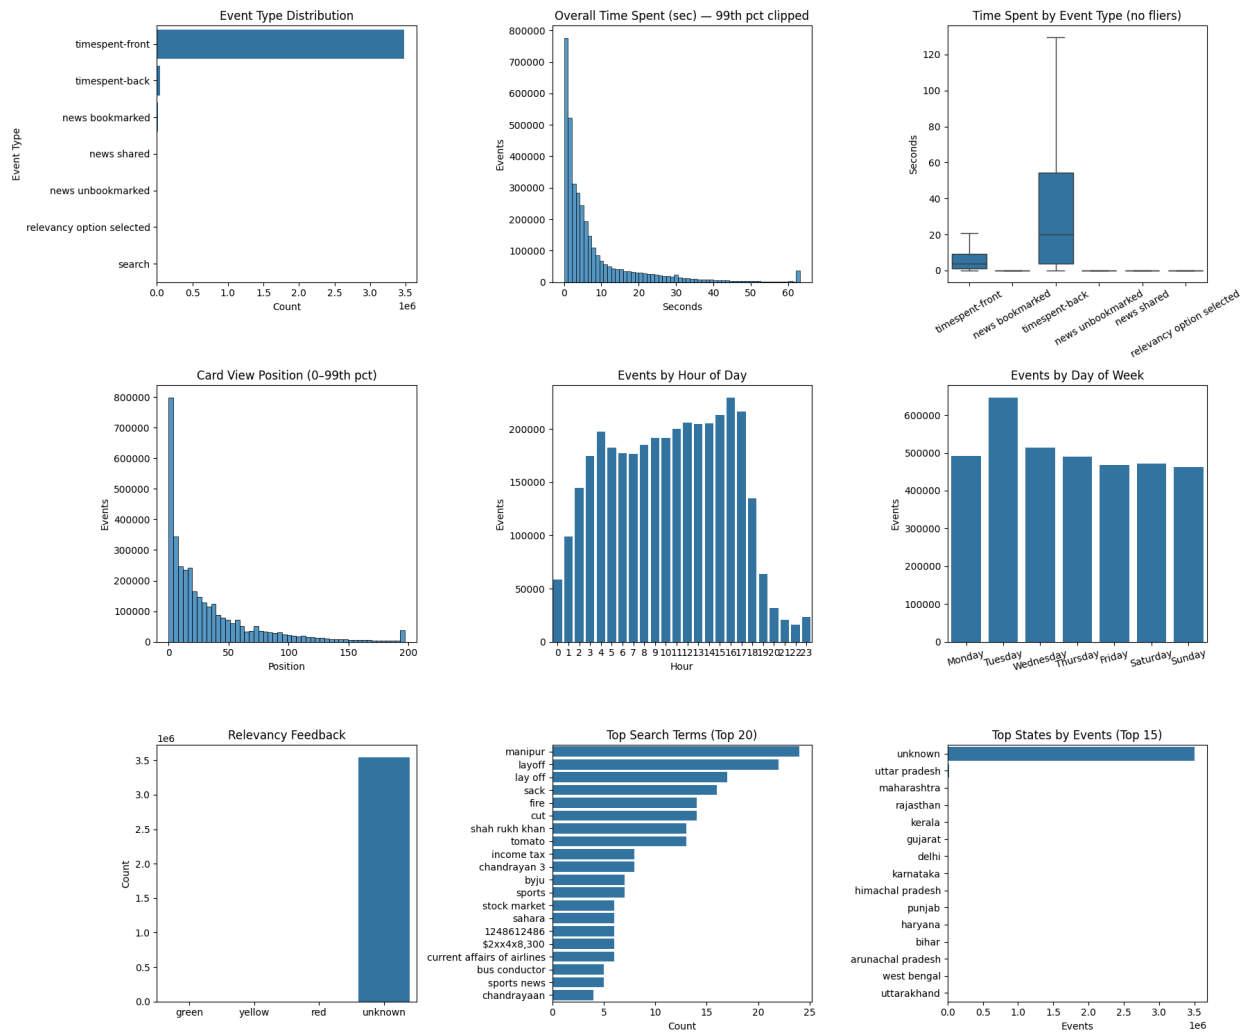
- ◆ Search Terms

- Top searches: manipur, layoff, fire, Shah Rukh Khan, income tax, Chandrayaan 3, stock market, sports.

◆ Geographic Distribution

- Large portion marked as “unknown” location (~3.5M).
- Among known: Uttar Pradesh, Maharashtra, Rajasthan, Kerala, Gujarat, Delhi dominate.

Events EDA Dashboard



- **Algorithms**

Both approaches rely on **embedding-based retrieval** using FAISS. The main difference lies in the **search space** during validation.

Events table is **split chronologically** using `eventTimestamp`:

- **Train Set:** 80% (older events).
- **Validation Set:** 20% (most recent events).

- ◆ **Approach 1 – Full Content Embeddings**

- **Embedding Space:** All content (train + validation).
- **User Representation:**
 - For each user, select **Top-K news interacted with**, based on interaction score.
 - Compute **weighted average embedding** of these Top-K items.
- **Validation Process:**
 - For each user in the validation set, retrieve **nearest neighbours from all available content**.
 - Re-rank candidates using **recency**.

- ◆ **Approach 2 – Validation-Only Embeddings**

- **Embedding Space:** Only validation content.
- **User Representation:** Same as Approach 1 (Top-K weighted average).
- **Validation Process:**
 - Nearest neighbours are retrieved only from the **validation set** (i.e., news available at that time).
 - Re-ranked with **recency**.

3. Algorithm Components

Data Preprocessing

- **Custom CSV Parser** developed to handle **multilingual content** and **emoticons**, ensuring correct tokenization and embedding input.

Interaction Scoring

- Each user's **Top-K news** are chosen based on a **composite score**:
 - **Event Type**: Higher weight for strong signals (e.g., bookmarks > clicks).
 - **Time Spent**: Longer reading time indicates higher interest.
 - **Relevancy Color Feedback**: Explicit signals (green > yellow > red).

Cold Start Strategy

- For new users with no history:
 - Recommend Top-K **popular and recent** news articles.
- Ensures recommendations even for first-time users.

Embedding Generation

- **Multilingual model** `multilingual-e5-base` used to cover diverse language inputs.
- News embedding = **weighted average** of:
 - Title embedding.
 - Content embedding.
 - Metadata embedding.

Vector Indexing

- Embeddings are **L2-normalized**.
Stored in **FAISS index** for efficient nearest neighbour retrieval.
- Supports large-scale retrieval in milliseconds.

Reranking


- After FAISS retrieval, candidate news are **re-ranked based on recency**.

Experiment Tracking:

- All runs logged with **MLflow** (embeddings, metrics, parameters).

● Results

Orange is approach 1 and Blue is approach 2



5. Productionization Plan

User-Specific Workflow

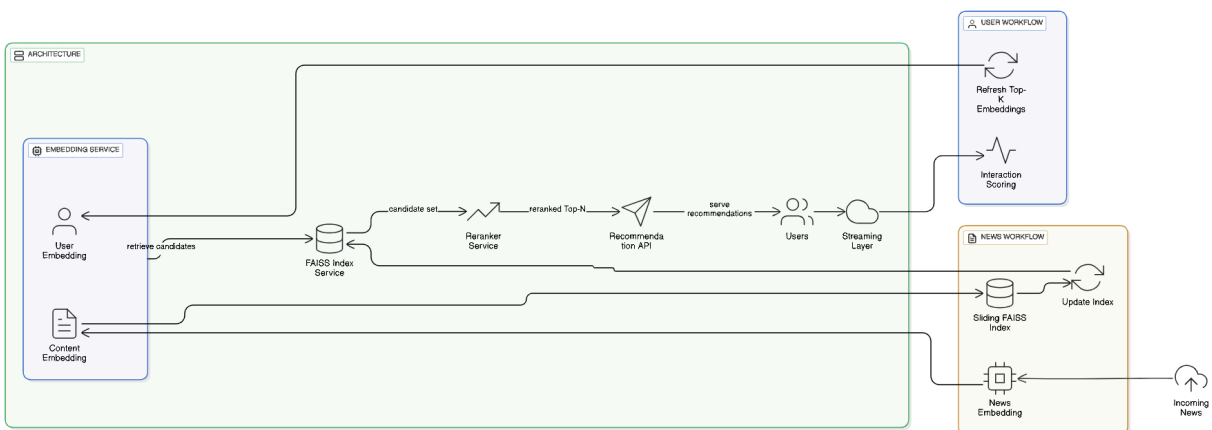
- Generate **interaction score** in real time whenever a user engages with news.
- Refresh **Top-K embeddings per user** every hour (or custom interval).
- Update **recommendations every 12 hours** with latest content.

News-Specific Workflow

- Generate **embeddings for incoming news in real time**.
- Maintain a **sliding FAISS index** of last 1 month's news.
- Update index hourly (or custom interval) to ensure freshness.


Architecture at Scale

1. **Streaming Layer:** Collect user interactions
2. **Embedding Service:** Generate/update embeddings (user + content).
3. **FAISS Index Service:** Maintain and query nearest neighbour index.
4. **Reranker Service:** Re-rank retrieved candidates with recency + popularity.
5. **Recommendation API:** Serve Top-N news to users.



- **Ending notes:**

- I have considered python to be equivalent to SQL queries and it's not like I am not confident with SQL. its just I didn't spend much time to think on that
- Code requires atleast 20 GB so running on atleast 32 GB RAM laptop is preferable
- Generating embedding is the most resource and time consuming task so atleast 12 core is preferable to optimize runtime



- I have thought of many future directions to optimize the solution but I don't have much time so leaving it as it is
- ChatGPT proved to be a very helpful tool in completing this assignment. I have used ChatGPT for following things:
 - Understanding domain
 - Generating boilerplate code